




Deployment Guide for Research made Readable

This guide provides instructions for deploying the Research made Readable application using multiple deployment methods:

1.  **Docker Deployment** (Recommended - Easiest)
2.  **Traditional VM Deployment** (Manual installation)
3.  **Cloud Platform Deployment** (Platform-specific)

Prerequisites

For Docker Deployment (Recommended)

- Docker and Docker Compose installed on the host system
- 2GB RAM minimum (4GB recommended)
- 5GB disk space minimum
- Your AbacusAI API key

For Traditional VM Deployment

- Ubuntu 20.04 LTS or later
- Python 3.8 or higher
- 2GB RAM minimum (4GB recommended)
- 10GB disk space minimum

Simplified Deployment: No external database installation required! The application uses DuckDB with Parquet file storage for a completely self-contained setup.



Docker Deployment (Recommended)

Docker deployment provides the easiest and most reliable way to deploy the Research made Readable application. This method ensures consistent behavior across different operating systems and eliminates dependency issues.

Docker Deployment Benefits

- ✓ **Consistent Environment:** Same behavior across all platforms
- ✓ **Easy Setup:** No manual dependency management
- ✓ **Isolation:** Application runs in its own container
- ✓ **Data Persistence:** Automatic data backup and persistence
- ✓ **Easy Updates:** Simple container rebuild process
- ✓ **Scalability:** Easy to scale horizontally

Docker Installation

On Ubuntu/Debian

```
# Update package index
sudo apt update

# Install Docker
sudo apt install -y docker.io docker-compose

# Add user to docker group
sudo usermod -aG docker $USER

# Start and enable Docker service
sudo systemctl start docker
sudo systemctl enable docker

# Log out and log back in for group changes to take effect
```

On CentOS/RHEL

```
# Install Docker
sudo yum install -y docker docker-compose

# Start and enable Docker service
sudo systemctl start docker
sudo systemctl enable docker

# Add user to docker group
sudo usermod -aG docker $USER
```

On Windows/macOS

Download and install [Docker Desktop](https://www.docker.com/products/docker-desktop/) (<https://www.docker.com/products/docker-desktop/>)

Docker Deployment Steps

1. Clone the application repository

```
git clone <repository-url> research_summary_app
cd research_summary_app
```

1. Set up environment variables

```
# Copy the Docker environment example file
cp .env.docker-example .env

# Edit the .env file with your API key
nano .env # or use your preferred text editor
```

1. Configure your API key

```
# In the .env file, replace the placeholder with your actual API key
ABACUSAI_API_KEY=your_actual_abacusai_api_key_here
```

1. Deploy with Docker Compose

```
# Build and start the application
docker-compose up -d --build

# Check if the container is running
docker-compose ps

# View logs
docker-compose logs -f
```

1. Access the application

- Open your browser and navigate to `http://your-server-ip:8501`
- For local deployment: `http://localhost:8501`

Docker Management Commands

```
# Start the application
docker-compose up -d

# Stop the application
docker-compose down

# Restart the application
docker-compose restart

# View logs
docker-compose logs -f research-app

# Update the application
git pull origin main
docker-compose down
docker-compose up -d --build

# Backup data
tar -czf research_backup_$(date +%Y%m%d_%H%M%S).tar.gz data/

# Restore data
tar -xzf research_backup_YYYYMMDD_HHMMSS.tar.gz
```

Docker Production Configuration

For production environments, consider these additional configurations:

1. Use a reverse proxy (nginx)

```
# docker-compose.production.yml
version: '3.8'
services:
  nginx:
    image: nginx:alpine
    ports:
      - "80:80"
      - "443:443"
    volumes:
      - ./nginx.conf:/etc/nginx/nginx.conf
      - ./ssl:/etc/nginx/ssl
    depends_on:
      - research-app
    restart: unless-stopped

  research-app:
    build: .
    expose:
      - "8501"
    environment:
      - ABACUSAI_API_KEY=${ABACUSAI_API_KEY}
    volumes:
      - ./data:/app/data
      - ./logs:/app/logs
    restart: unless-stopped
```

1. Set up SSL/TLS certificates

```
# Use Let's Encrypt for free SSL certificates
sudo apt install certbot
sudo certbot --nginx -d yourdomain.com
```

1. Configure monitoring and health checks

```
# Add to docker-compose.yml
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost:8501/_stcore/health"]
  interval: 30s
  timeout: 10s
  retries: 3
  start_period: 40s
```

Docker Troubleshooting

Common Issues and Solutions

Port Already in Use

```
# Check what's using port 8501
sudo netstat -tulnp | grep :8501

# Kill the process using the port
sudo kill -9 <process_id>

# Or use a different port in docker-compose.yml
ports:
  - "8502:8501"
```

Permission Issues

```
# Fix data directory permissions
sudo chown -R $USER:$USER ./data ./logs
chmod -R 755 ./data ./logs
```

Container Won't Start

```
# Check container logs
docker-compose logs research-app

# Check if all required environment variables are set
docker-compose config

# Verify Docker daemon is running
sudo systemctl status docker
```

API Key Issues

```
# Verify .env file exists and is properly configured
cat .env

# Check container environment variables
docker-compose exec research-app env | grep ABACUSAI

# Test API key manually
docker-compose exec research-app python -c "
import os
print('API Key:', os.getenv('ABACUSAI_API_KEY', 'NOT SET'))
"
```

Database/Data Issues

```
# Reset database (WARNING: This will delete all data)
docker-compose down
rm -rf ./data/db/*
docker-compose up -d

# Check data directory permissions
ls -la ./data/

# Verify data persistence
docker-compose down
docker-compose up -d
# Your data should still be there
```

Memory Issues

```
# Check container memory usage
docker stats research-made-readable

# Increase memory limit in docker-compose.yml
services:
  research-app:
    mem_limit: 2g
    mem_reservation: 1g
```

Docker Security Best Practices

1. Environment Variables

```
# Use Docker secrets for sensitive data
echo "your_api_key_here" | docker secret create abacusai_api_key -

# Reference in docker-compose.yml
secrets:
  - abacusai_api_key
```

1. Network Security

```
# Use custom networks
networks:
  research-network:
    driver: bridge
    ipam:
      config:
        - subnet: 172.20.0.0/16
```

1. User Permissions

```
# Run as non-root user (already configured in Dockerfile)
USER streamlit
```

1. Resource Limits

```
# Set resource limits
deploy:
  resources:
    limits:
      cpus: '2.0'
      memory: 2G
    reservations:
      memory: 1G
```



Traditional VM Deployment

For users who prefer manual installation or have specific requirements that need customization.

VM Setup

1. Initial Server Setup

```
# Update system packages
sudo apt update && sudo apt upgrade -y

# Install required system packages
sudo apt install -y python3 python3-pip python3-venv git nginx supervisor

# Create application user
sudo useradd -m -s /bin/bash research-made-readable
sudo usermod -aG sudo research-made-readable
```

2. Database Setup

No Database Setup Required! The application uses DuckDB with Parquet files for data storage. Database files are created automatically in the `data/db/` directory when the application first runs.

Benefits of this approach:

- No external database server installation or configuration
- Automatic database initialization on first startup
- Portable data storage - simply copy the `data/db/` directory to backup or migrate
- No database credentials or connection strings to manage

3. Application Deployment

```
# Switch to application user
sudo -u research-made-readable -i

# Clone the application
git clone <repository-url> /home/research-made-readable/research_summary_app
cd /home/research-made-readable/research_summary_app

# Create virtual environment
python3 -m venv venv
source venv/bin/activate

# Install dependencies
pip install -r requirements.txt

# Set up environment configuration (see detailed section below)
cp .env-example .env
nano .env # Edit with your actual API key

# Set secure permissions
chmod 600 .env

# Create required directories
mkdir -p data/db data/uploads data/exports logs

# Initialize application (database files created automatically)
python setup.py
```

4. Environment Configuration

API Key Setup (REQUIRED)

The application requires an AbacusAI API key to access all AI models. Here's how to set it up securely in production:

1. Obtain your AbacusAI API key:

- Visit [AbacusAI](https://abacus.ai/) (<https://abacus.ai/>)
- Sign up for an account or log in
- Navigate to your account settings or API section
- Generate a new API key
- Copy the key for configuration

2. Configure environment variables:

```
```bash
Edit the .env file with your actual API key
nano .env
```

# Add your API key (replace with your actual key)

```
ABACUSAI_API_KEY=your_actual_api_key_here
```

# Optional: Add other configuration settings

```
LOG_LEVEL=INFO
```

```
DEBUG=False
```

```
MAX_UPLOAD_SIZE=10
```

```
```
```

1. Secure the environment file:

```
```bash
Set restrictive permissions (owner read/write only)
chmod 600 .env
```

# Verify permissions

```
ls -la .env
```

# Should show: -rw----- 1 research-made-readable research-made-readable

```
```
```

Production Security Considerations

Environment Variables Security:

- Never commit the `.env` file to version control
- Use restrictive file permissions (600 or 640 maximum)
- Consider using system environment variables for enhanced security:

```
bash
```

```
# Alternative: Set system environment variable
```

```
echo 'export ABACUSAI_API_KEY="your_key_here"' >> ~/.bashrc
```

```
source ~/.bashrc
```

API Key Management:

- Use separate API keys for development and production
- Monitor API usage through your AbacusAI dashboard
- Set up usage alerts to prevent unexpected charges
- Rotate API keys regularly for security

Network Security:

- Restrict API access to your server's IP if possible
- Use HTTPS in production (see SSL Certificate section)
- Consider using a reverse proxy for additional security

**Cost Management and Monitoring****API Usage Optimization:**

- Monitor API usage through AbacusAI dashboard
- Set up billing alerts and usage limits
- Use lower-cost models (e.g., GPT-4-Mini) for development/testing
- Implement rate limiting to prevent abuse

Usage Tracking:

```
# Monitor application logs for API usage
tail -f logs/streamlit.log | grep -i "api\|error"

# Check supervisor logs
tail -f /var/log/research-made-readable.log
```

**Troubleshooting Environment Issues****Common API Key Issues:****1. Invalid API Key Error:**

```
```bash
Check if API key is set
grep ABACUSAI_API_KEY .env

Test API key validity
curl -H "Authorization: Bearer your_api_key_here" \
https://apps.abacus.ai/v1/chat/completions
```
```

1. Permission Errors:

```
bash
# Fix file permissions
sudo chown research-made-readable:research-made-readable .env
chmod 600 .env
```

2. Environment Not Loading:

```
```bash
Verify python-dotenv is installed
pip list | grep python-dotenv

Check if .env file exists in correct location
ls -la .env
```
```

Environment Validation Script:

```
# Create a simple validation script
cat > validate_env.py << 'EOF'
import os
from dotenv import load_dotenv

load_dotenv()

api_key = os.getenv('ABACUSAI_API_KEY')
if api_key:
    print(f"✅ API key found: {api_key[:10]}...")
else:
    print("❌ API key not found")

print(f"✅ Working directory: {os.getcwd()}")
print(f"✅ .env file exists: {os.path.exists('.env')}")
EOF

python validate_env.py
```

5. Nginx Configuration

```
# Create Nginx configuration
sudo nano /etc/nginx/sites-available/research-made-readable

# Add configuration:
server {
    listen 80;
    server_name your_domain.com; # Replace with your domain

    location / {
        proxy_pass http://localhost:8501;
        proxy_http_version 1.1;
        proxy_set_header Upgrade $http_upgrade;
        proxy_set_header Connection 'upgrade';
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
        proxy_cache_bypass $http_upgrade;
        proxy_read_timeout 86400;
    }
}

# Enable site
sudo ln -s /etc/nginx/sites-available/research-made-readable /etc/nginx/sites-enabled/
sudo nginx -t
sudo systemctl reload nginx
```

6. Supervisor Configuration

```
# Create supervisor configuration
sudo nano /etc/supervisor/conf.d/research-made-readable.conf

# Add configuration:
[program:research-made-readable]
command=/home/research-made-readable/research_summary_app/venv/bin/streamlit run
app.py --server.port=8501 --server.address=localhost
directory=/home/research-made-readable/research_summary_app
user=research-made-readable
autostart=true
autorestart=true
redirect_stderr=true
stdout_logfile=/var/log/research-made-readable.log
environment=PATH="/home/research-made-readable/research_summary_app/venv/bin"

# Update supervisor
sudo supervisorctl reread
sudo supervisorctl update
sudo supervisorctl start research-made-readable
```

7. SSL Certificate (Optional but Recommended)

```
# Install Certbot
sudo apt install certbot python3-certbot-nginx

# Get SSL certificate
sudo certbot --nginx -d your_domain.com

# Auto-renewal
sudo crontab -e
# Add: 0 12 * * * /usr/bin/certbot renew --quiet
```

Monitoring and Maintenance

1. Application Monitoring

```
# Check application status
sudo supervisorctl status research-made-readable

# View logs
sudo tail -f /var/log/research-made-readable.log

# Restart application
sudo supervisorctl restart research-made-readable
```

2. Data Backup and Maintenance

```
# Backup all data including environment configuration
cd /home/research-made-readable/research_summary_app

# Create comprehensive backup (data + configuration)
tar -czf backup_$(date +%Y%m%d).tar.gz data/db/ logs/ .env-example
# Note: .env file is excluded for security - document API key separately

# Alternative: Copy data directory only
cp -r data/db/ ../backups/backup_$(date +%Y%m%d)/

# Restore data (simple file copy)
cd /home/research-made-readable/research_summary_app
tar -xzf backup_20240101.tar.gz

# Verify data integrity (optional)
python -c "
import duckdb
conn = duckdb.connect('data/db/research_app.duckdb')
print('Tables:', conn.execute('SHOW TABLES').fetchall())
conn.close()
print('Data verification complete')
"
```

Environment Configuration Backup:

- **Do NOT backup the .env file** in version control or regular backups
- **Securely document your API key** in a password manager or secure note
- **Keep a copy of .env-example** for reference when setting up new environments
- **Test environment restoration** by validating API key access after restore

Automated Backup Script:

```
#!/bin/bash
# Create automated backup script
cat > backup_app.sh << 'EOF'
#!/bin/bash
BACKUP_DIR="/home/research-made-readable/backups"
APP_DIR="/home/research-made-readable/research_summary_app"
DATE=$(date +%Y%m%d_%H%M%S)

mkdir -p $BACKUP_DIR
cd $APP_DIR

# Backup data and logs (excluding sensitive .env file)
tar -czf $BACKUP_DIR/research_app_backup_$DATE.tar.gz \
    data/db/ \
    logs/ \
    .env-example \
    --exclude='.env'

echo "Backup completed: $BACKUP_DIR/research_app_backup_$DATE.tar.gz"

# Keep only last 7 days of backups
find $BACKUP_DIR -name "research_app_backup_*.tar.gz" -mtime +7 -delete
EOF

chmod +x backup_app.sh

# Add to crontab for daily backup
crontab -e
# Add: 0 2 * * * /home/research-made-readable/research_summary_app/backup_app.sh
```

3. System Updates

```
# Update system packages
sudo apt update && sudo apt upgrade -y

# Update Python packages
cd /home/research-made-readable/research_summary_app
source venv/bin/activate
pip install --upgrade -r requirements.txt

# Restart services
sudo supervisorctl restart research-made-readable
```

Security Considerations

1. Firewall Configuration

```
# Install UFW
sudo apt install ufw

# Configure firewall
sudo ufw default deny incoming
sudo ufw default allow outgoing
sudo ufw allow ssh
sudo ufw allow 'Nginx Full'
sudo ufw enable
```

2. Application Security

```
# Set proper file permissions
sudo chown -R research-made-readable:research-made-readable /home/research-made-readable/research_summary_app
sudo chmod -R 755 /home/research-made-readable/research_summary_app
sudo chmod 600 /home/research-made-readable/research_summary_app/.env

# Configure secure headers in Nginx
sudo nano /etc/nginx/sites-available/research-made-readable
# Add to server block:
add_header X-Frame-Options DENY;
add_header X-Content-Type-Options nosniff;
add_header X-XSS-Protection "1; mode=block";
add_header Strict-Transport-Security "max-age=31536000; includeSubDomains";
```

3. Data Security

```
# Secure data directory permissions
sudo chown -R research-made-readable:research-made-readable /home/research-made-readable/research_summary_app/data/
sudo chmod -R 750 /home/research-made-readable/research_summary_app/data/db/

# Secure environment file
sudo chmod 600 /home/research-made-readable/research_summary_app/.env

# Note: DuckDB files are local to the application - no network security concerns
```

Backup and Recovery

1. Automated Backups

```
# Create backup script
sudo nano /home/research-made-readable/backup.sh

#!/bin/bash
DATE=$(date +%Y%m%d_%H%M%S)
BACKUP_DIR="/home/research-made-readable/backups"
APP_DIR="/home/research-made-readable/research_summary_app"
mkdir -p $BACKUP_DIR

# Data backup (DuckDB and Parquet files)
tar -czf $BACKUP_DIR/data_backup_$DATE.tar.gz -C $APP_DIR data/db/

# Full application backup (including data)
tar -czf $BACKUP_DIR/app_backup_$DATE.tar.gz $APP_DIR

# Clean old backups (keep last 30 days)
find $BACKUP_DIR -name "*_backup_*.tar.gz" -mtime +30 -delete

echo "Backup completed: $DATE"

# Make executable
sudo chmod +x /home/research-made-readable/backup.sh

# Add to crontab for daily backups
sudo crontab -e
# Add: 0 2 * * * /home/research-made-readable/backup.sh
```

2. Recovery Procedures

```
# Stop application first
sudo supervisorctl stop research-made-readable

# Restore data only (DuckDB and Parquet files)
cd /home/research-made-readable/research_summary_app
tar -xzf /home/research-made-readable/backups/data_backup_YYYYMMDD_HHMMSS.tar.gz

# OR restore entire application
tar -xzf /home/research-made-readable/backups/app_backup_YYYYMMDD_HHMMSS.tar.gz -C /
home/research-made-readable/

# Set proper permissions
sudo chown -R research-made-readable:research-made-readable /home/research-made-readable/research_summary_app
sudo chmod -R 750 /home/research-made-readable/research_summary_app/data/db/

# Restart application
sudo supervisorctl start research-made-readable
```

Performance Optimization

1. Data Storage Optimization

```
# DuckDB automatically optimizes performance, but you can:

# Monitor disk space usage
df -h /home/research-made-readable/research_summary_app/data/db/

# Check Parquet file sizes
ls -lah /home/research-made-readable/research_summary_app/data/db/*.parquet

# Ensure sufficient disk space for data growth
# DuckDB compresses data efficiently with Parquet format
# Typical compression ratio: 5-10x compared to raw CSV data

# For large datasets, consider SSD storage for better I/O performance
```

2. Application Optimization

```
# Configure Streamlit for production
nano /home/research-made-readable/research_summary_app/.streamlit/config.toml

[server]
maxUploadSize = 200
maxMessageSize = 200
enableCORS = false
enableXsrfProtection = true

[browser]
gatherUsageStats = false

[theme]
primaryColor = "#2563EB"
backgroundColor = "#FFFFFF"
secondaryBackgroundColor = "#F8F9FA"
textColor = "#1F2937"
```

Troubleshooting

Common Issues

1. Application won't start

- Check supervisor logs: `sudo tail -f /var/log/researchlens.log`
- Verify environment variables in `.env`
- Check database connectivity

2. Data storage issues

- Check if `data/db/` directory exists and is writable
- Verify Parquet files are not corrupted: `python -c "import duckdb; duckdb.connect('data/db/research_app.duckdb').execute('SHOW TABLES')"`
- Check disk space: `df -h /home/research-made-readable/research_summary_app/data/db/`

3. Nginx errors

- Check Nginx configuration: `sudo nginx -t`
- Review Nginx logs: `sudo tail -f /var/log/nginx/error.log`

4. SSL certificate issues

- Check certificate status: `sudo certbot certificates`
- Renew certificate: `sudo certbot renew`

Performance Issues

1. Slow data operations

- DuckDB automatically optimizes queries
- Check available disk space for temporary operations
- Consider SSD storage for better I/O performance
- Monitor Parquet file sizes for unexpected growth

2. High memory usage

- Monitor with `htop` or `top`
- Adjust Streamlit configuration
- DuckDB has efficient memory management for Parquet files
- Consider upgrading server resources for large datasets

Maintenance Schedule

Daily

- Monitor application logs
- Check system resources
- Verify backup completion
- Monitor disk space in `data/db/` directory

Weekly

- Review security logs
- Update system packages
- Check Parquet file sizes and growth trends
- Verify data directory permissions

Monthly

- Analyze application performance
- Review DuckDB storage efficiency
- Update application dependencies
- Security audit
- Test backup and recovery procedures



DuckDB/Parquet Architecture Benefits

Simplified Deployment

- **No Database Server:** Eliminated PostgreSQL installation and configuration
- **Self-Contained:** All data stored in portable Parquet files
- **Reduced Dependencies:** Fewer system packages and services to manage

- **Faster Setup:** Database initialization happens automatically

Operational Advantages

- **Simple Backups:** Copy the `data/db/` directory - that’s it!
- **Easy Migration:** Transfer the entire application directory to any server
- **No Database Credentials:** No connection strings or passwords to manage
- **Portable Development:** Identical setup for development and production

Performance & Reliability

- **Optimized Storage:** Parquet format provides excellent compression and query performance
- **ACID Compliance:** DuckDB ensures data integrity
- **Automatic Optimization:** No manual database tuning required
- **Efficient Memory Usage:** DuckDB designed for analytical workloads

Deployment Comparison

| Aspect | PostgreSQL (Before) | DuckDB + Parquet (After) |
|-----------------|--|--------------------------|
| System Packages | 8+ packages including PostgreSQL | 5 basic packages |
| Database Setup | Manual configuration, users, permissions | Automatic initialization |
| Backup Method | <code>pg_dump</code> + application files | Simple directory copy |
| Migration | Database dump/restore + files | Copy entire directory |
| Security | Database passwords, network config | File permissions only |
| Dependencies | External database service | Self-contained |

This deployment guide provides a comprehensive setup for production deployment of Research made Readable on an external VM with the new simplified DuckDB/Parquet architecture. The deployment process is now significantly simpler and more reliable than the previous PostgreSQL-based setup.